

Yonatan Erez

HW4

Ps.c:

```
C ps.c > ...
1  #include "types.h"
2  #include "stat.h"
3  #include "user.h"
4  #include "fcntl.h"
5
6  int main(int argc, char const *argv[])
7  {
8      ps();
9      exit();
10 }
11
12 |
```

Creates a terminal interface to the user.

Demosched.c:

```
C demosched.c > main(int, char const * [])
1  #include "types.h"
2  #include "stat.h"
3  #include "user.h"
4  #include "fcntl.h"
5
6  int main(int argc, char const *argv[])
7  {
8      getprio();
9      ps();
10     long num = 1 << 28;
11     int pid = fork();
12     if (pid == -1) {
13         printf(1, "\nError in fork\n");
14         exit();
15     }
16     if (pid == 0) {
17         setprio(7);
18         ps();
19         for (volatile long l1=0; l1<num; l1++){
20             l1++;
21             l1--;
22         }
23         printf(1, "\nChild exit\n");
24         exit();
25     }
26     else {
27         setprio(6);
28         ps();
29         for (volatile long l2=0; l2<num; l2++){
30             l2++;
31             l2--;
32         }
33         printf(1, "\nParent is waiting for child\n");
34         wait();
35     }
36     ps();
37     printf(1, "Parent exit\n");
38     exit();
39 }
40
```

The script uses `getprio`, `setprio` and `ps` (for debug purposes).

```

init: starting sh
$ demosched
o
o getprio: pid=3, priority=0

pid      state      priority    ticks_counter  required_ticks
1        SLEEPING      0           5               1
2        SLEEPING      0           3               1
3        RUNNING      0           1               1
o
o setprio: pid=3, priority=6
o pid      state      priority    ticks_counter  required_ticks
o 1        SLEEPING      0           5               1
  2        SLEEPING      0           3               1
  3        RUNNING      6           2              64
  4        RUNNABLE     0           0               1

setprio: pid=4, priority=7

pid      state      priority    ticks_counter  required_ticks
1        SLEEPING      0           5               1
2        SLEEPING      0           3               1
3        RUNNABLE     6          128             64
4        RUNNING      7           0             128

Parent is waiting for child

Child exit

```

```

pid      state      priority    ticks_counter  required_ticks
1        SLEEPING      0           5               1
2        SLEEPING      0           3               1
3        RUNNING      6          128             64
Parent exit

```

Let's add another ps call.

```

pid      state      priority    ticks_counter  required_ticks
1        SLEEPING      0           5               1
2        SLEEPING      0           3               1
5        RUNNING      0           2               1
$ 

```

As we can see the sequence is:

1. Boot.
2. Demosched.
3. Getprio is called by the parent.
4. Fork (new process is added to that table).
5. Setprio(6) is called by the parent.
6. Setprio(6) is called by the child.
7. Parent is running till the required ticks are over.
8. Parent finished calculations and waits for child.
9. Child is running till the required ticks are over.
10. Child finished calculations and exits.
11. Parent exits.
12. Ps is called from the terminal and both parent and child do not exist.

Makefile:

```
181     _usertests\  
182     _wc\  
183     _zombie\  
184     _demosched\  
185     _ps\  
186  
187     fs.img: mkfs README $(UPROGS)  
188     ./mkfs fs.img README $(UPROGS)
```

```
251  
252     EXTRA=\  
253     mkfs.c ulib.c user.h cat.c echo.c forktest.c grep.c kill.c\  
254     ln.c ls.c mkdir.c rm.c stressfs.c usertests.c wc.c zombie.c\  
255     printf.c umalloc.c demosched.c ps.c\  
256     README dot-bochsrc *.pl toc.* runoff runoff1 runoff.list\  
257     .gdbinit.tmpl gdbutil\  
258
```

Both PS and Makefile were added to the makefile compilation lines.

Defs.h:

```
121 void      wakeup(void);  
122 void      yield(void);  
123 int       proc_ps(void);  
124 int       proc_getprio(int);  
125 int       proc_setprio(int, int);  
126
```

The new syscalls implementation are decalred.

Proc.c:

```
545
546 int
547 v proc_getprio(int pid)
548 {
549     int priority = -1;
550     struct proc* p;
551     acquire(&ptable.lock);
552 v for (p=ptable.proc; p<&ptable.proc[NPROC]; p++){
553 v     if (p->pid == pid){
554         priority = p->priority;
555         cprintf("\n");
556         cprintf("getprio: pid=%d, priority=%d\n", p->pid, p->priority);
557         break;
558     }
559 }
560 release(&ptable.lock);
561 return priority;
562 }
563
```

```
566
567 int
568 v proc_setprio(int pid, int priority)
569 {
570     struct proc* p;
571     acquire(&ptable.lock);
572 v for (p=ptable.proc; p<&ptable.proc[NPROC]; p++){
573 v     if (p->pid == pid){
574         p->priority = priority;
575         cprintf("\n");
576         cprintf("setprio: pid=%d, priority=%d\n", p->pid, p->priority);
577         break;
578     }
579 }
580 release(&ptable.lock);
581 return 0;
582 }
583
```



```

587 int
588 proc_ps()
589 {
590     struct proc* p;
591     acquire(&table.lock);
592     cprintf("\n");
593     cprintf("pid \t state \t          priority \t ticks_counter \t  required_ticks\n");
594     for (p=table.proc; p<&table.proc[NPROC]; p++){
595         if (p->state == UNUSED){
596             continue;
597         }
598         if (p->state == EMBRYO){
599             cprintf("%d \t EMBRYO \t %d \t \t %d \t          %d \n", p->pid, p->priority, p->ticks_counter, 1 << p->priority);
600             continue;
601         }
602         if (p->state == SLEEPING){
603             cprintf("%d \t SLEEPING \t %d \t \t %d \t          %d \n", p->pid, p->priority, p->ticks_counter, 1 << p->priority);
604             continue;
605         }
606         if (p->state == RUNNABLE){
607             cprintf("%d \t RUNNABLE \t %d \t \t %d \t          %d \n", p->pid, p->priority, p->ticks_counter, 1 << p->priority);
608             continue;
609         }
610         if (p->state == RUNNING){
611             cprintf("%d \t RUNNING \t %d \t \t %d \t          %d \n", p->pid, p->priority, p->ticks_counter, 1 << p->priority);
612             continue;
613         }
614         if (p->state == ZOMBIE){
615             cprintf("%d \t ZOMBIE \t %d \t \t %d \t          %d \n", p->pid, p->priority, p->ticks_counter, 1 << p->priority);
616             continue;
617         }
618     }
619     release(&table.lock);
620     return 0;
621 }

```

The implementations are added to proc.c file.

```

342     // before jumping back to us.
343     c->proc = p;
344     switchuvm(p);
345     p->state = RUNNING;
346     swtch(&(c->scheduler), p->context);
347     volatile int* tick_counter_ptr = &p->ticks_counter;
348     while (*tick_counter_ptr < (1 << p->priority))
349     {
350         if (p->state != RUNNING){
351             break;
352         }
353         continue;
354     }
355
356     switchkvm();
357
358     // Process is done running for now.
359     // It should have changed its p->state before coming back.
360     c->proc = 0;
361 }
362 release(&table.lock);
363
364 }
365 }
366

```

The scheduler picks the lowest pid which did not finish ticking.

Proc.h:

```
45 struct trapframe *t; // trap frame for current syscall
46 struct context *context; // switch() here to run process
47 void *chan; // If non-zero, sleeping on chan
48 int killed; // If non-zero, have been killed
49 struct file *ofile[NOFILE]; // Open files
50 struct inode *cwd; // Current directory
51 char name[16]; // Process name (debugging)
52 int priority; // Process priority
53 int ticks_counter; // Process ticks counter
54 };
55
```

The new attributes are added to the proc struct.

Syscall.c:

```
104 extern int sys_write(void);
105 extern int sys_uptime(void);
106 extern int sys_getprio(void);
107 extern int sys_setprio(void);
108 extern int sys_ps(void);
109
128 [SYS_unlink] sys_unlink,
129 [SYS_link] sys_link,
130 [SYS_mkdir] sys_mkdir,
131 [SYS_close] sys_close,
132 [SYS_getprio] sys_getprio,
133 [SYS_setprio] sys_setprio,
134 [SYS_ps] sys_ps
135 };
136
137 void
138 syscall(void)
```

Added the new syscalls to the list.

Proc.h:

```
19 #define SYS_unlink 18
20 #define SYS_link 19
21 #define SYS_mkdir 20
22 #define SYS_close 21
23 #define SYS_getprio 22
24 #define SYS_setprio 23
25 #define SYS_ps 24
26
```

The syscalls number are added.

Sysproc.c:

```
93 |  
94 | int  
95 | ✓ sys_getprio(void){  
96 | | return proc_getprio(myproc()->pid);  
97 | }  
98 |  
99 |  
100 | int  
101 | ✓ sys_setprio(void){  
102 | | int priority;  
103 | | ✓ if (argint(0, &priority) < 0){  
104 | | | return -1;  
105 | | }  
106 | | return proc_setprio(myproc()->pid, priority);  
107 | }  
108 |  
109 |  
110 | int  
111 | ✓ sys_ps(void){  
112 | | return proc_ps();  
113 | }
```

The implementations are being called by the syscalls.

Trap.c:

```
102
103 // Force process to give up CPU on clock tick.
104 // If interrupts were on while locks held, would need to check nlock.
105 if(myproc() && myproc()->state == RUNNING &&
106    tf->trapno == T_IRQ0+IRQ_TIMER){
107     myproc()->ticks_counter++;
108     if (myproc()->ticks_counter > (1<<7)){
109         myproc()->ticks_counter = 1<<7;
110         yield();
111     }
112 }
113
114 // Check if the process has been killed since we yielded
115 if(myproc() && myproc()->killed && (tf->cs&3) == DPL_USER)
116     exit();
117
118
```

On every iteration the ticks_counter is incremented till max value of $2^7 = 128$.

If the counter reaches the *counter_ticks* = $2^{priority}$.

User.h:

```
21 int dup(int);
22 int getpid(void);
23 char* sbrk(int);
24 int sleep(int);
25 int uptime(void);
26 int getprio(void);
27 int setprio(int);
28 int ps(void);
29
30 // ulib.c
31 int stat(const char*, struct stat*);
```

The new system calls are added to the user space.

Usys.S:

```
30 SYSCALL(sleep)
31 SYSCALL(uptime)
32 SYSCALL(getprio)
33 SYSCALL(setprio)
34 SYSCALL(ps)
35 |
```

The syscalls are added to the list.